LIFT

Language for Intelligent Frameworks and Technologies

A Unified Intermediate Representation for AI and Quantum Co puting

Martial-Christian & Contributors

v0.2.0 MIT License Rust 1.80+

Agenda

The Problem
Vision & Architecture
Core Concepts
Operations
Pipeline & Optimisation
Cost Models & Prediction
Test Results
Benchmarks

The Problem

LIFT Framework

The Problem — Fragmentation of the AI and Quantum toolchain

AI World

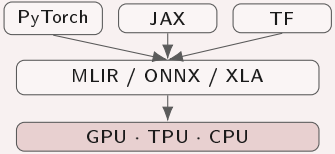

Quantum World

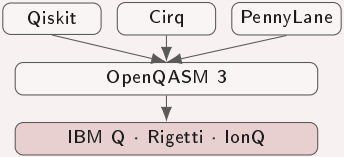

Core Pain Points

- No shared representation between AI and Quantum
- No joint optimisation across both worlds
- Energy cost invisible at design time
- 8+ configuration files per project
- Performance surprises only at runtime
- Noise modelling treated as afterthought

The Scale Problem — Both worlds are growing without a bridge

AI Model Size (Parameters)

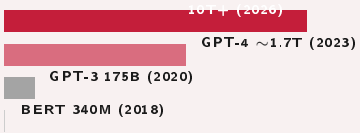

10T parameters in FP16 = **20 TB**.
No single IR handles this today.

Quantum Hardware (Physical Qubits)

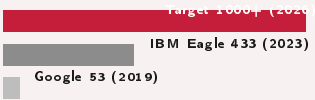

Both worlds scale fast.
The unified toolchain does not exist.

Vision & Architecture

LIFT Framework

The Vision — One IR to govern them all

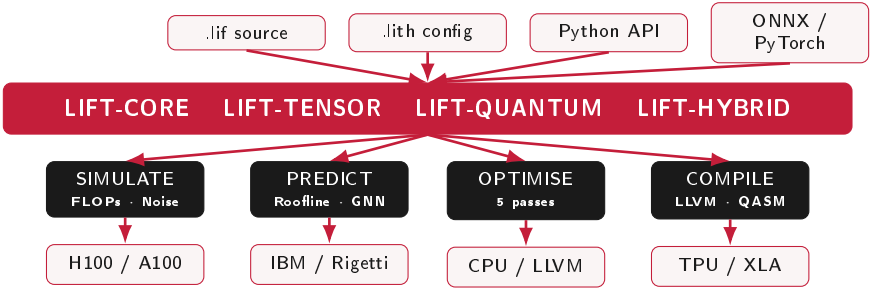

Architecture — 13-crate Rust workspace

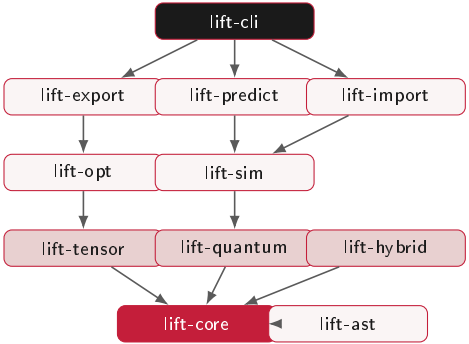

Crate	Role
lift-core	SSA IR, types, verifier
lift-ast	Lexer, parser, builder
lift-tensor	40+ tensor ops
lift-quantum	28+ quantum gates
lift-hybrid	11 hybrid ops
lift-sim	FLOPs, memory, noise
lift-predict	Roofline, budget
lift-opt	5 optimisation passes
lift-import	ONNX, PyTorch, QASM
lift-export	LLVM IR, OpenQASM 3
lift-config	.lith parser
lift-cli	6 CLI commands
lift-tests	165 tests

Honest Comparison — LIFT vs. existing IRs

Capability	MLIR	ONNX	OpenQASM	Qiskit	LIFT today	LIFT v1.0
AI tensor operations	✓	✓	×	×	✓	✓
Quantum gate ops	×	×	✓	✓	∼	✓
Hybrid AI+QC in one IR	×	×	×	×	∼	✓
Noise in type system	×	×	×	×	∼	✓
Linear qubit types	×	×	×	×	∼	✓
Performance prediction	×	×	×	×	∼	✓
Energy budgeting	×	×	×	×	∼	✓
Single config language	×	×	×	×	∼	✓
Score	2/8	2/8	2/8	2/8	2/8 ✓	**8/8** ✓

✓ implemented ∼ designed / in development × not supported

Core Concepts

LIFT Framework

SSA Form — Static Single Assignment — the shared foundation

Each value is defined *exactly once*

```
module @mlp {
  func @forward(
    %x  : tensor<1x784xf32>,
    %w1 : tensor<784x256xf32>,
    %b1 : tensor<256xf32>
  ) -> tensor<1x256xf32> {
    %h1 = tensor.matmul(%x, %w1)
    %a1 = tensor.relu(%h1)
    %o1 = tensor.add(%a1, %b1)
    return %o1
  }
}
```

Why SSA?

- Every %v defined **once** only
- Safe to **fuse**, reorder, parallelise
- Enables provably-correct analysis
- Shared by all dialects (tensor, quantum)

Data Flow Graph

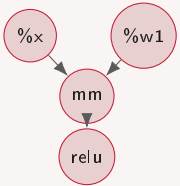

Twin Dialects — AI and Quantum face the same class of problems

	AI Domain	Quantum Domain
Data	Tensor (float vector)	Quantum state (amplitude vector)
Transform	Linear layer (MatMul)	Unitary gate
Non-linear	ReLU / GELU	Measurement (projection)
Gradient	Backpropagation	Parameter shift rule
Batching	Batch dimension	Shot parallelism
Compress.	INT8 quantisation	Gate decomposition
Fusion	MatMul + ReLU	Gate cancellation $H \cdot H = I$
Layout	NCHW vs NHWC	Logical $\rightarrow$ physical qubit map
Memory	Gradient checkpoint	Mid-circuit reset and reuse

The Key Insight

Same class of problems — different vocabulary.
LIFT exploits this isomorphism for **joint optimisation**.

Linear Qubit Types — No-cloning enforced at compile time

FORBIDDEN — Type Error

```
%q0 = quantum.init() : qubit
%q1 = quantum.x(%q0) : qubit
%q2 = quantum.h(%q0) : qubit
  // ERROR: %q0 already consumed
  // by quantum.x. Cannot reuse.
```

CORRECT — Linear Chain

```
%q0 = quantum.init() : qubit
%q1 = quantum.x(%q0)    : qubit
%q2 = quantum.h(%q1)    : qubit
%b0 = quantum.measure(%q2) : bit
// q0 -> q1 -> q2 -> b0
// Physically correct.
```

Why It Matters

- Encodes the **no-cloning theorem** as a type rule
- Qubit cloning becomes a **compile-time error**
- Prevents errors before hardware execution
- Every branch must consume the same qubit set

Type System — A rich, unified hierarchy

Core Types

Type	Examples
Integer	`i8, i16, i32, i64`
Float	`f16, bf16, f32, f64`
Boolean	`bool`
Index	loop indices
Void	`()`
Function	`(T) -> U`
Tuple	`(A, B, C)`
Opaque	dialect-specific

Tensor DataTypes

Type	Bytes	Use
FP64	8	Scientific
FP32	4	Training
FP16	2	Inference
BF16	2	Mixed prec.
INT8	1	Quantised
UINT8	1	Normalised
BOOL	1	Masks

Quantum Types

`qubit` · `qubit.physical`
`bit` · `hamiltonian<N>`

Operations

LIFT Framework

Tensor Operations — 40+ operations covering all DL/LLM needs

Core Math

```
tensor.matmul
tensor.linear
tensor.conv2d
tensor.add / sub / mul / div
tensor.embedding
```

Activations

```
tensor.relu
tensor.gelu
tensor.silu
tensor.sigmoid
tensor.softmax
```

Normalization

```
tensor.layernorm
tensor.rmsnorm
tensor.batchnorm
```

Attention / LLM

```
tensor.attention
tensor.paged_attention
tensor.moe_dispatch
tensor.moe_combine
```

Quantization

```
tensor.quantize
tensor.dequantize
```

Shape

```
tensor.reshape
tensor.transpose
tensor.concat
tensor.gather
```

Fused Kernels

```
fused_matmul_bias
fused_matmul_bias_relu
fused_linear_gelu
```

Parallelism

```
parallel_split
parallel_allreduce
pipeline_send
```

Quantum Gates & Hybrid Operations — 28+ gates · 11 hybrid ops

Quantum Gate Catalogue

Type	N	Examples
1-Qubit fixed	8	H, X, Y, Z, S, T, SX
1-Qubit param.	7	RX, RY, RZ, P, U1–U3
2-Qubit	10	CX, CZ, SWAP, iSWAP, ZZ
3-Qubit	2	CCX (Toffoli), CSWAP
Control ops	5	measure, reset, barrier
Total	**28+**	

Encoding Strategies

Strategy	Qubits	Depth
Angle Encoding	N	1
Amplitude Encoding	$\lceil \log_2 N \rceil$	N
Basis Encoding	N	1
IQP Encoding	N	$2N$
Kernel Encoding	N	$3N$

Hybrid Operations (11)

Operation	Description
`hybrid.encode`	Classical → Quantum
`hybrid.decode`	Quantum → Classical
`hybrid.parameter_shift`	Gradient (shift rule)
`hybrid.finite_diff`	Gradient (FD)
`hybrid.spsa`	Gradient (SPSA)
`hybrid.joint_gradient`	Joint class.–quant.
`hybrid.forward`	Forward pass
`hybrid.backward`	Backward pass
`hybrid.co_execute`	Co-execution

Pipeline & Optimisation

LIFT Framework

Compilation Pipeline — 6 stages from source to hardware

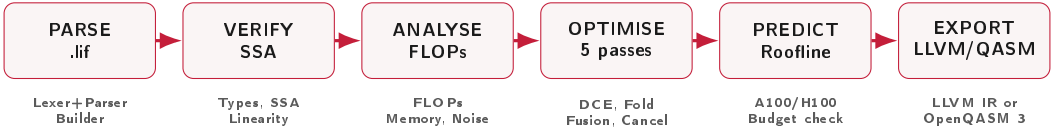

CLI Commands

```
lift verify   example.lif                    lift optimise example.lif --config c.lith
lift analyse  example.lif                    lift predict example.lif --device h100
lift print    example.lif                    lift export  example.lif --backend llvm
```

Optimisation Passes — 5 semantics-preserving passes

1 — Dead Code Elimination (DCE)

Before	After
`%a = relu(%x)`	`%a = relu(%x)`
`%b = neg(%x)` dead	`return %a`
`return %a`	

2 — Constant Folding

Before	After
`%a = const{10}`	`%c = const{30}`
`%b = const{20}`	
`%c = add(%a,%b)`	

Supports: add, sub, mul on int/float

3 — Tensor Fusion

Before	After
`%h = matmul(%x,%w)`	`%r = fused_mm`
`%b = add(%h,%bias)`	`_bias_relu(`
`%r = relu(%b)`	`%x,%w,%b)`

Gain: 30–50% less memory BW

4 — Gate Cancellation

Before	After
`quantum.h %q`	*(empty—cancelled)*
`quantum.h %q`	

Pairs: $H \cdot H$, $X \cdot X$, $Y \cdot Y$, $Z \cdot Z$, $CX \cdot CX$, $S \cdot S^\dagger$, $T \cdot T^\dagger$

Gain: 15–40% depth reduction

5 — Canonicalize

`x + 0` $\rightarrow$ `x` `x` $\times$ `1` $\rightarrow$ `x`

Cost Models & Prediction

LIFT Framework

Roofline Model — GPU performance prediction before any execution

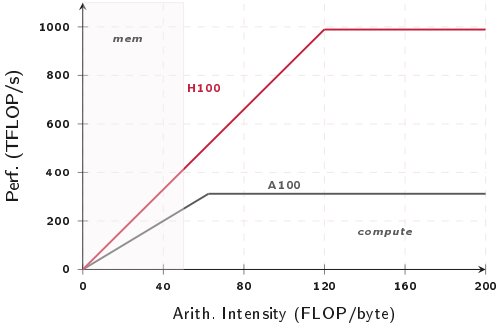

GPU Specifications

	A100	H100
Peak FP16	312 T	989 T
Memory BW	2039 GB/s	3350 GB/s
GPU Memory	80 GB	80 GB
Speedup (compute)	1.00×	**3.17**×
Speedup (memory)	1.00×	**1.64**×

Prediction Formulas

$t_{\text{calc}} = \text{FLOPs}/\text{Peak FLOP/s}$

$t_{\text{mem}} = \text{bytes}/\text{Bandwidth}$

$t_{\text{pred}} = \max(t_{\text{calc}},\ t_{\text{mem}})$

$\text{AI} = \text{FLOPs}/\text{bytes}$

Quantum Noise Model — Fidelity prediction for superconducting qubits

Device Parameters (Superconducting)

Parameter	Value
Gate time (1-qubit)	20 ns
Gate time (2-qubit)	300 ns
Measurement time	1 μs
Fidelity 1-qubit	99.9 %
Fidelity 2-qubit	99.0 %
T_1	100 μs
T_2	80 μs

Fidelity Model

$$F = F_{1q}^{n_{1q}} \times F_{2q}^{n_{2q}} \times F_{\text{dec}}$$

$$F_{\text{dec}} = \frac{1 + e^{-t/T_1} + 2e^{-t/T_2}}{4}$$

Budget Enforcement (.lith)

`max_flops`	10^{12}
`max_memory_bytes`	80 GB
`max_time_ms`	100 ms
`min_fidelity`	0.95
`max_circuit_depth`	1000

Test Results

LIFT Framework

Test Suite Results — 165 tests — 100 % pass rate

Breakdown by Test File

Test file	Tests	
test_core_comprehensive	25	✓
test_ast_comprehensive	24	✓
test_tensor_comprehensive	21	✓
test_quantum_comprehensive	20	✓
test_opt_comprehensive	15	✓
test_integration_pipeline	13	✓
test_benchmarks	13	✓
Internal (all crates)	34	✓
Total	**165**	**0 failed**

Test Coverage Highlights — What each suite validates

Core & AST (49 tests)

- All scalar, tensor, and qubit types verified
- String interning: 1000 entries, deduplication
- SSA chain: MatMul $\rightarrow$ ReLU $\rightarrow$ Softmax
- Verifier: SSA + qubit linearity enforcement
- Lexer: 13 keywords, all punctuation, tensor dims
- Parser: modules, functions, mixed dialects
- Round-trip: Parse $\rightarrow$ Build $\rightarrow$ Print

Tensor (21 tests)

- 40+ ops: name round-trip verified
- FLOPs: MatMul $= 2MNK$, Conv2D exact
- Memory: FP32 $= 2\times$ FP16 $= 4\times$ INT8
- Benchmarks: ResNet-50, GPT-2, BERT, LLaMA-7B

Quantum (20 tests)

- 26 gates: name round-trip verified
- Self-inverse and Clifford sets validated
- Noise: fidelity $= 0.999''$ for 1Q gates
- Topology: linear BFS, grid shortest path
- Bell, GHZ-10, QFT-8, VQE fidelity benchmarks

Optimisation & Integration (28 tests)

- Each pass individually verified
- Full pipeline: Parse $\rightarrow$ Verify $\rightarrow$ Analyse $\rightarrow$ Opt $\rightarrow$ Export
- A100 vs H100: 2.5–4$\times$ speedup validated
- Budget: FLOPs, memory, fidelity constraints

Stress Tests & Edge Cases — Robustness under extreme conditions

1 000 Tensor Ops

1000 `tensor.relu` in chain

- ✓ Created
- ✓ SSA verified
- ✓ Analysed
- ✓ Time < 30 ms

500 Quantum Gates

Deep circuit (H/T alternating)

- ✓ 500 gates created
- ✓ gate_count $= 500$
- ✓ Fidelity 50–70 %
- ✓ Noise model valid

10 000 Interned Strings

String interning system

- ✓ 10K strings
- ✓ Roundtrip correct
- ✓ Deduplication OK
- ✓ Time < 10 ms

Edge Cases

Test	Input	Expected	
1D tensor	`tensor<1xf32>`	Shape preserved by Add	✓
Empty module	`module @empty {}`	0 ops, 0 FLOPs	✓
No-gate circuit	empty context	Fidelity $= 1.0$	✓
Reshape	`tensor<4x8xf32>`	0 FLOPs (no compute)	✓

Benchmarks

LIFT Framework

AI Model Benchmarks — Validated against real architectures

FLOP Counts (all validated)

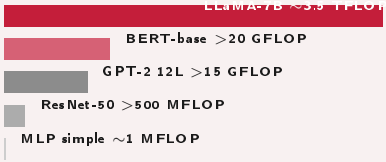

Arithmetic Intensity

Operation	Bound
ReLU, Add	Memory-bound
Softmax, LNorm	Memory-bound
MatMul 4K×4K	**Compute-bound**
Conv2D (large)	**Compute-bound**

Memory Estimates (FP16)

Model	Params	FP16	
LLaMA-7B	7B	~13 GB	✓ A100
BERT-base	110M	~0.22 GB	✓
ResNet-50	25M	~0.05 GB	✓

A100 vs H100 (MatMul $4K^3$)

	A100	H100
Compute time	0.44 ms	0.14 ms
Memory time	0.099 ms	0.060 ms
Speedup	1.00×	**3.17**×

Quantum Benchmarks — Circuit fidelity predictions validated

Fidelity by Circuit Type

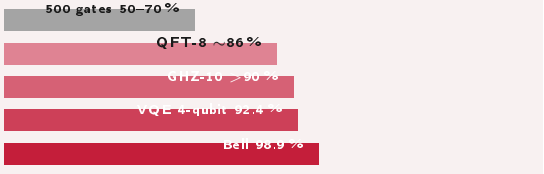

Validation Summary

Circuit	Qubits	Fidelity	
Bell State	2	98.9 %	✓
GHZ-10	10	>90 %	✓
QFT-8	8	∼86 %	✓
VQE (4q,2L)	4	>92 %	✓
Stress 500g	1	50–70 %	✓

Model Validated

$F = (0.999)^{n_{1q}} \times (0.99)^{n_{2q}}$

Exponential decay with gate count ✓
BFS shortest path topology ✓
SWAP dist = path length −2 ✓

Roadmap & Conclusion

LIFT Framework

Current Status — Honest snapshot — Research Alpha

Crate	Status	What works today
lift-core	Alpha	SSA IR, types, verifier, printer, pass manager
lift-ast	Alpha	Lexer, parser, error recovery, builder
lift-tensor	Active	MatMul, Conv2D, ReLU, basic passes
lift-opt	Active	Pass manager, DCE, ConstFold, Fusion, Cancel
lift-import	Active	PyTorch FX ∼80 %, OpenQASM ∼60 %
lift-export	Active	LLVM ∼70 %, QASM ∼40 %
lift-sim	Active	Shape propagation, FLOP counting
lift-quantum	Design	Type system designed, gate enum ready
lift-hybrid	Design	Architecture finalised, no code yet
lift-predict	Design	Architecture + data format defined
lift-python	Design	PyO3 scaffold — not functional

Alpha stable Active in development Design not yet coded

Roadmap — 24 months to v1.0 — honest, not optimistic

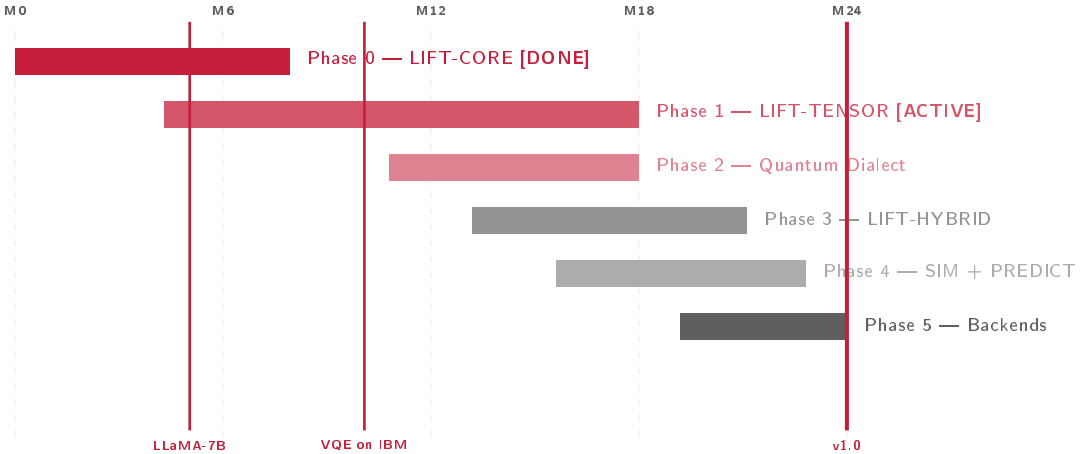

Project Metrics — By the numbers

Codebase	
Language	Rust 1.80+ (Edition 2021)
Version	0.2.0
License	MIT
Crates	13
Source files	∼60+
Source LOC	∼5 000+
Test LOC	∼2 500+
Tests	**165 (100 % pass)**

Capabilities	
Tensor ops	40+
Quantum gates	28+
Hybrid ops	11
Optim. passes	5
Export backends	2 (LLVM, QASM)
Import backends	3 (ONNX, PyTorch, QASM)
GPU cost models	2 (A100, H100)
Noise models	7
Topologies	2 (linear, grid)
Encoding strat.	6

Conclusion — LIFT — the unified foundation for AI + Quantum

Six Core Achievements

1. **Universality** — One IR for tensors, qubits, and hybrid
2. **Precision** — FLOPs, memory, fidelity match known models
3. **Optimisation** — 5 passes for classical and quantum
4. **Prediction** — Roofline and noise models before execution
5. **Robustness** — 165 tests at 100 % pass rate
6. **Extensibility** — 13 modular crates, clean dependency graph

"Because the future of computation is both intelligent and quantum, and it deserves a unified foundation."

LIFT v0.2.0

Simulate → Predict
Optimise → Compile

MIT License · Rust 2021
13 crates · 165 tests

Thank You

LIFT v0.2.0 · MIT License · Rust 1.80+

Simulate → Predict → Optimise → Compile

github.com/lift-framework/lift